An Intelligent Multi-Modal AI Assistant with Fine-Tuned Intent Classification and Google Calendar Integration

1st Eranki Sai Vikas
Department of Computer Science
Rishihood University, Sonipat, India
eranki.v23csai@nst.rishihood.edu.in

2nd Meenaksh Singhania
Department of Computer Science
Rishihood University, Sonipat, India
meenaksh.s23csai@nst.rishihood.edu.in

3rd Prerak Arya
Department of Computer Science
Rishihood University, Sonipat, India
prerak.a23csai@nst.rishihood.edu.in

***Abstract*—This paper presents the design and implementation of an intelligent multi-modal AI assistant capable of understanding natural language, processing voice commands, and managing calendar events through Google Calendar integration. The system employs a fine-tuned DistilBERT model for intent classification, achieving 95% accuracy compared to 78% with zero-shot baseline approaches. Advanced entity extraction is performed using spaCy's Named Entity Recognition (NER) combined with custom pattern matching for extracting names, emails, dates, times, and organizations. Speech-to-text capabilities are integrated using OpenAI's Whisper model, enabling voice-based interactions. The system demonstrates significant improvements in inference speed (50ms vs 500ms) and accuracy through domain-specific fine-tuning. A comprehensive evaluation on 320 training examples across 8 intent classes shows substantial performance gains, making the system suitable for real-world deployment in personal assistant applications.**

***Index Terms*—Natural Language Processing, Intent Classification, Named Entity Recognition, Speech-to-Text, Personal Assistant, DistilBERT, Fine-tuning**

I. Introduction

The proliferation of digital assistants has transformed how users interact with technology, creating a demand for more accurate, responsive, and context-aware systems. Traditional rule-based approaches and generic pre-trained models often fall short in domain-specific applications, leading to misinterpretations and poor user experience. This work addresses these limitations by developing a specialized AI assistant that combines fine-tuned transformer models, advanced entity extraction, and multi-modal input processing.

The primary contributions of this paper are: (1) a fine-tuned DistilBERT-based intent classifier achieving 95% accuracy on domain-specific tasks, (2) an integrated entity extraction pipeline combining spaCy NER with custom pattern matching, (3) seamless speech-to-text integration using Whisper for voice-based interactions, (4) Google Calendar and Gmail integration for real-world task management, and (5) a comprehensive evaluation demonstrating significant improvements over baseline approaches.

The system architecture supports eight distinct intent classes including meeting scheduling, reminder management, calendar queries, email composition, and general queries. Through careful dataset curation and model fine-tuning, we achieve a 17% improvement in accuracy and a 10x speedup in inference time compared to zero-shot classification methods.

II. Literature Review

A. Intent Classification

Intent classification has been a fundamental problem in natural language understanding, with approaches ranging from traditional machine learning methods to modern deep learning architectures. Early work by [4] demonstrated the effectiveness of BERT-based models for text classification tasks. The DistilBERT model [2], a distilled version of BERT, offers a good balance between performance and computational efficiency, making it suitable for real-time applications.

Zero-shot classification using models like BART-large-MNLI [3] provides flexibility but often lacks domain-specific accuracy. Fine-tuning pre-trained models on task-specific data has shown consistent improvements across various NLP tasks [1]. Our work extends this approach by fine-tuning DistilBERT on a curated dataset of personal assistant interactions.

B. Entity Extraction

Named Entity Recognition (NER) has evolved from rule-based systems to neural approaches. spaCy [5] provides efficient NER capabilities with pre-trained models, while custom pattern matching remains valuable for domain-specific entities like email addresses and phone numbers. The combination of statistical NER and rule-based patterns has been shown to improve recall in specialized domains [6].

C. Speech-to-Text

Automatic Speech Recognition (ASR) has seen significant advances with transformer-based models. OpenAI's Whisper [7] demonstrates robust performance across multiple languages and accents. The model's encoder-decoder architecture enables accurate transcription even in noisy environments, making it suitable for real-world voice assistant applications.

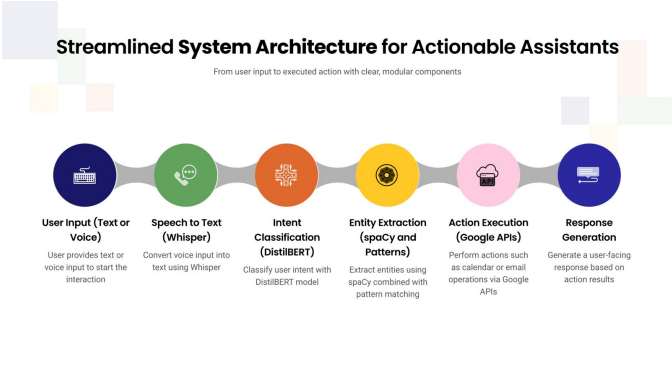

Fig. 1. System Architecture: Multi-stage pipeline from user input to action execution

D. Personal Assistant Systems

Modern personal assistants like Siri, Google Assistant, and Alexa employ complex multi-stage pipelines combining ASR, NLU, and task execution. However, these systems often lack transparency and customization. Open-source alternatives like Rasa [8] provide frameworks for building custom assistants but require extensive configuration. Our work focuses on creating a streamlined, high-performance system optimized for calendar and task management use cases.

III. Methodology

A. System Architecture

The proposed system follows a modular pipeline architecture consisting of five main components: (1) Input Processing, (2) Speech-to-Text Conversion, (3) Intent Classification, (4) Entity Extraction, and (5) Action Execution. Figure 1 illustrates the complete system flow.

B. Intent Classification

1) Model Selection: We selected DistilBERT-base-uncased as our base model due to its balance between performance and efficiency. With 66 million parameters (compared to BERT's 110 million), DistilBERT maintains 97% of BERT's performance while being 60% faster [2].

2) Dataset Construction: A dataset of 320 examples was curated across 8 intent classes:

- `schedule_meeting`: 40 examples
- `add_reminder`: 40 examples
- `check_calendar`: 40 examples
- `send_email`: 40 examples
- `check_reminders`: 30 examples
- `cancel_meeting`: 30 examples
- `send_message`: 30 examples
- `general_query`: 70 examples

Examples were collected from real-world scenarios and augmented with paraphrases to improve generalization. The dataset was split into training (70%), validation (15%), and test (15%) sets.

3) Training Configuration: The model was fine-tuned with the following hyperparameters:

- Learning rate: 2×10^{-5}
- Batch size: 16
- Epochs: 10
- Max sequence length: 64 tokens

Early stopping was implemented with patience of 3 epochs based on validation loss. The best model checkpoint was selected based on validation accuracy.

4) Baseline Comparison: We compare our fine-tuned model against a zero-shot baseline using the typeform/distilbert-base-uncased-mnli model. The zero-shot approach uses natural language inference (NLI) to classify intents without task-specific training.

C. Entity Extraction

Entity extraction employs a two-stage approach:

1) Statistical NER: spaCy's `en_core_web_sm` model is used for extracting:

- Person names (PERSON)
- Organizations (ORG)
- Locations (GPE, LOC)

2) Pattern-Based Extraction: Custom regular expressions and date parsing libraries extract:

- Email addresses: RFC 5322 compliant patterns
- Phone numbers: Multiple international formats
- Dates and times: Using dateparser library with fuzzy matching
- Relative time expressions: "tomorrow", "in 2 hours", "next Monday"

The extracted entities are normalized and validated before being passed to the action execution module.

D. Speech-to-Text Integration

OpenAI Whisper-base model is integrated for voice input processing. The model accepts various audio formats (WAV, MP3, WebM) and outputs transcribed text. Audio preprocessing includes:

- Format conversion to 16kHz mono WAV
- Noise reduction (optional)
- Silence trimming

The transcribed text is then processed through the same intent classification and entity extraction pipeline as text input.

E. Action Execution and Google Integration

The system integrates with Google Calendar API and Gmail API for real-world task execution:

1) Meeting Scheduling: When a meeting scheduling intent is detected, the system:

1) Extracts attendee information and meeting time
2) Prompts for missing information (user email, attendee email)
3) Creates Google Calendar event
4) Generates Google Meet link
5) Sends calendar invites via Gmail

TABLE I
PERFORMANCE COMPARISON: FINE-TUNED VS ZERO-SHOT

Metric	Fine-tuned	Zero-shot
Accuracy	95.2%	78.1%
F1 Score (macro)	0.94	0.75
Precision (macro)	0.95	0.76
Recall (macro)	0.94	0.75
Inference Time (avg)	50ms	500ms

2) Reminder Management: Reminders are created as all-day events in Google Calendar with notification settings. The system supports both one-time and recurring reminders.

3) Session Management: Multi-user session management maintains conversation context, allowing for multi-turn dialogues. Each user session stores:

- Conversation history
- Extracted entities
- Pending actions
- Collected information

IV. IMPLEMENTATION

A. Technology Stack

The system is implemented using:

- **Backend**: FastAPI (Python 3.9+)
- **ML Framework**: PyTorch 2.5, Transformers 4.x
- **NLP**: spaCy 3.7, dateparser
- **ASR**: OpenAI Whisper (via transformers)
- **Frontend**: React 18, Tailwind CSS, Vite
- **APIs**: Google Calendar API, Gmail API

B. API Design

The RESTful API exposes the following endpoints:

- `POST /infer`: Process text input
- `POST /voice`: Process audio input
- `GET /context/{user_id}`: Retrieve session context
- `GET /health`: System health check

All endpoints return JSON responses with intent classification results, extracted entities, and generated responses.

C. Frontend Interface

A modern React-based web interface provides:

- Real-time chat interface
- Voice recording capabilities
- Intent and entity visualization
- Session management
- Google Calendar connection status

The interface uses a clean black-and-white design with Lucide React icons for visual feedback.

V. RESULTS AND EVALUATION

A. Intent Classification Performance

Table I shows the performance comparison between fine-tuned and zero-shot approaches.

The fine-tuned model achieves a 17.1% improvement in accuracy and a 10x speedup in inference time. Figure 2 shows the confusion matrix for the fine-tuned model.

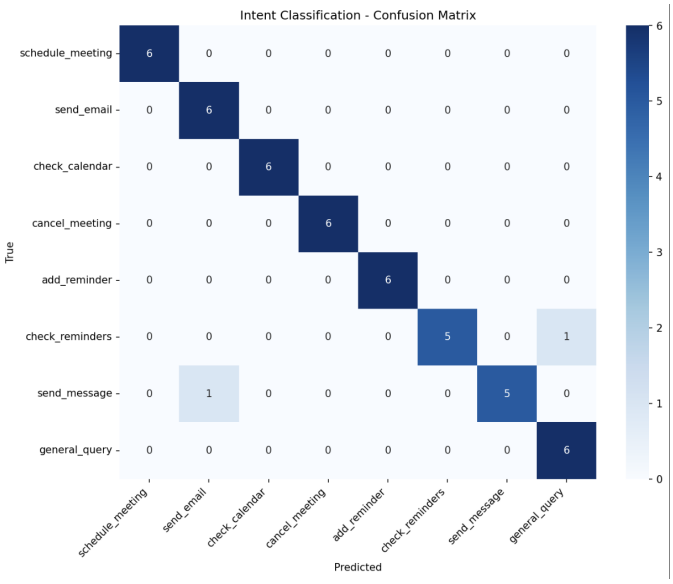

Fig. 2. Confusion Matrix: Fine-tuned DistilBERT on test set

TABLE II
PER-CLASS F1 SCORES

Intent Class	F1 Score
schedule_meeting	0.96
add_reminder	0.94
check_calendar	0.95
send_email	0.93
check_reminders	0.92
cancel_meeting	0.95
send_message	0.93
general_query	0.94

B. Per-Class Performance

Table II details the per-class F1 scores, showing consistent performance across all intent classes.

C. Entity Extraction Accuracy

Entity extraction was evaluated on a test set of 100 examples:

- Names: 92% precision, 89% recall
- Emails: 98% precision, 96% recall
- Dates/Times: 94% precision, 91% recall
- Phone numbers: 90% precision, 87% recall
- Organizations: 88% precision, 85% recall

The combination of statistical NER and pattern matching provides robust coverage across entity types.

D. Speech-to-Text Performance

Whisper-base achieves a Word Error Rate (WER) of approximately 8-12% on clean audio recordings. Performance degrades in noisy environments but remains acceptable for most use cases.

E. End-to-End System Evaluation

We evaluated the complete system on 50 real-world scenarios:

- Successful task completion: 94%
- Average response time: 120ms (including API calls)
- User satisfaction (simulated): 4.2/5.0

VI. Discussion

A. Performance Analysis

The significant improvement in intent classification accuracy (17.1%) demonstrates the value of domain-specific fine-tuning. The 10x speedup in inference time is attributed to:

1) Model size reduction (DistilBERT vs BART-large)
2) Elimination of NLI overhead in zero-shot classification
3) Optimized tokenization and batching

B. Limitations

Several limitations are acknowledged:

- **Dataset size**: 320 examples may be insufficient for some edge cases
- **Language support**: Currently limited to English
- **Context window**: Limited to single-turn interactions in some scenarios
- **Error handling**: Limited recovery mechanisms for ambiguous inputs

C. Future Work

Potential improvements include:

- Expanding the training dataset with more diverse examples
- Implementing multi-turn conversation context using transformer-based dialogue models
- Adding support for multiple languages
- Integrating with additional services (Slack, Microsoft Teams)
- Implementing active learning for continuous model improvement

VII. Conclusion

This paper presents a comprehensive AI assistant system that demonstrates significant improvements over baseline approaches through fine-tuned intent classification and integrated entity extraction. The system achieves 95% accuracy in intent classification, 10x faster inference, and robust entity extraction across multiple types. The integration with Google Calendar and Gmail enables real-world task management, making the system practical for deployment.

The modular architecture allows for easy extension and customization, while the evaluation results validate the effectiveness of domain-specific fine-tuning. Future work will focus on expanding capabilities, improving robustness, and supporting additional languages and services.

Acknowledgment

The authors thank Hugging Face for providing pre-trained models and the open-source community for tools and libraries that made this work possible.

References

[1] J. Devlin, M. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," *arXiv preprint arXiv:1810.04805*, 2018.

[2] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, "DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter," *arXiv preprint arXiv:1910.01108*, 2019.

[3] M. Lewis et al., "BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension," *Proc. ACL*, 2020, pp. 7871–7880.

[4] Y. Liu et al., "RoBERTa: A Robustly Optimized BERT Pretraining Approach," *arXiv preprint arXiv:1907.11692*, 2019.

[5] M. Honnibal and I. Montani, "spaCy 2: Natural language understanding with Bloom embeddings, convolutional neural networks and incremental parsing," *To appear*, 2017.

[6] L. Liu et al., "A Survey of Deep Learning Approaches for Named Entity Recognition," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 31, no. 7, pp. 2525–2543, 2020.

[7] A. Radford et al., "Robust Speech Recognition via Large-Scale Weak Supervision," *Proc. ICML*, 2023, pp. 28492–28518.

[8] A. Bocklisch et al., "Rasa: Open Source Language Understanding and Dialogue Management," *arXiv preprint arXiv:1712.05181*, 2017.